

Introduction to Artificial Intelligence

DA 221

Jan - May 2023

IIT Guwahati


Instructors: Neeraj Sharma & Arghyadip Roy

Lecture 2

Tentative

AI History

- 1940 - 1950: Early days
- 1950 - 1970: Excitement and expectations
- 1970 - 1990: Knowledge-based approaches
- 1990 - 2005: Statistical approaches
- 2005 - Present: Rise of Neural Networks



Let's start with getting aware of some AI developments

- Time
- Researchers
- Problems
- Places and Companies

Artificial Intelligence


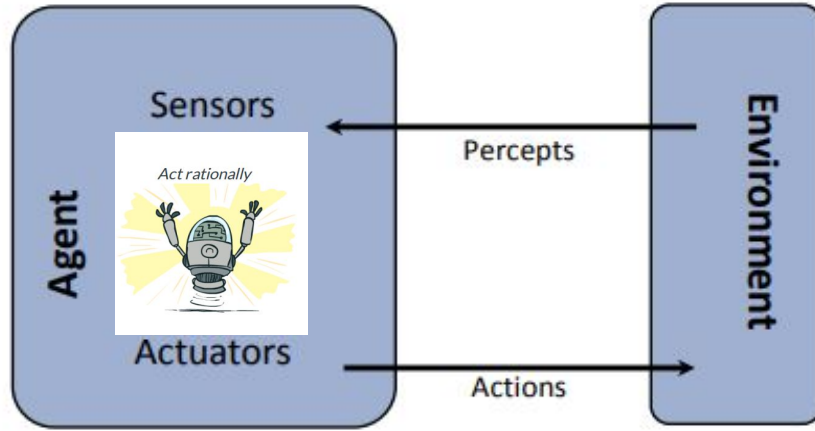



Image credits: CS188, UC Berkeley


Agents

An agent is an entity that perceives its environment through sensors and take actions through actuators.




Some examples

Pacman




Pacman

Percept sequence	Action
(left cell, no food)	go right
(left cell, food)	eat
(right cell, no food)	go left
(left cell, food)	eat
(left cell, no food), (left cell, no food)	go right
(left cell, no food), (left cell, food)	eat
(...)	(...)




Pacman

- The optimal Pacman?
- What is the right agent function?
- How to formulate the goal of Pacman?
 - 1 point per food dot collected up to time ?
 - 1 point per food dot collected up to time , minus one per move?
 - Penalize when too many food dots are left not collected?
- Can it be implemented in a small and efficient agent program?



Rational Agent

- Informally, a **rational agent** is an agent that does the "right thing".
- A **performance measure** evaluates a sequence of environment states caused by the agent's behavior.
- A rational agent is an agent that chooses whichever action that **maximizes** the **expected** value of the performance measure, given the percept sequence to date.



Rational Agent

- Rationality is **not equal** to omniscience
 - percepts may not supply all relevant information
- Rationality **is not equal** to clairvoyance action
 - outcomes may not be as expected
- Rational **is not equal** to successful

Rationality leads to exploration, learning and autonomy.

Some tasks

Example 1: a self-driving car

- **performance measure**: safety, destination, legality, comfort, ...
- **environment**: streets, highways, traffic, pedestrians, weather, ...
- **actuators**: steering, accelerator, brake, horn, speaker, display, ...
- **sensors**: video, accelerometers, gauges, engine sensors, GPS, ...

Some tasks

Example 2: an Internet shopping agent

- **performance measure:** price, quality, appropriateness, efficiency
- **environment:** current and future WWW sites, vendors, shippers
- **actuators:** display to user, follow URL, fill in form, ...
- **sensors:** web pages (text, graphics, scripts)

Let's define a problem



Let's define a problem



A cab driver driving through traffic

A problem

- Define a problem space
 - Are there sub-problems?
- Task


A problem

- Goal
- Constraints
- Task

A problem

- Goal
- Constraints
- Task

PEAS: Performance, Environment, Actuators, and Sensors



Lets' focus on types of Environment

- Fully observable vs. partially observable
- Single agent vs. multi-agent
- Deterministic vs. stochastic
- Episodic vs. sequential
- Static vs. dynamic
- Discrete vs. continuous
- Known vs unknown

Types of Environment

- Fully observable vs. partially observable
 - Single agent vs. multi-agent
 - Deterministic vs. stochastic
 - Episodic vs. sequential
 - Static vs. dynamic
 - Discrete vs. continuous
 - Known vs unknown
- Crossword puzzle
 - Chess, with a clock
 - Poker
 - Backgammon
 - Taxi driving
 - Medical diagnosis
 - Image analysis
 - Part-picking robot
 - Refinery controller
 - The real world

Thank you